# Reaktor

# Predictive analytics: hands-on session

Dev day 4.10.2013
Helsinki, Finland

# What is predictive analytics?

- **Analytics**
  - Actionable information from observations
  - Summaries, causality, predictions, graphs, ...
- **Predictive Analytics**
  - Predict future, unseen, or results of actions
  - *Models*, often probabilistic, are fitted to data
  - Statistics, machine learning, data mining

# Predictive analytics in practice

- Problem definition with business understanding
- Data collection
- Data preparation and exploration
- Modeling and evaluation
- Deployment
- Real world applications:
  - Customer relationship management: Netflix
  - Health care: Patients at risk of getting disease
  - Finance: Risk and fraud detection



Reaktor



# The Wine data

- Chemical measurements of variants of the Portuguese "Vinho Verde" wine
- For 1599 red and 4898 white wines
- **Objective**: Predict the color of wine based on the chemical measurements (see right panel)
- Notice: No data about grapes, brand, selling price etc. available

- Fixed acidity
- Volatile acidity
- Citric acid
- Residual sugar
- Chlorides
- Free sulfur dioxide
- Total sulfur dioxide
- Density
- pH
- Sulphates
- Alcohol



Reaktor

# Know your problem and data

- It is essential
  - to understand the domain
  - to get familiar with data
    - identify possible quality problems
    - discover possibilities and limitations
    - summary statistics and visualization
- Helps making correct choices in modeling
  - test various models

# Insight to the data

- To make models, you need to know your data.
- Summary statistics
  - Mean (average)
  - Standard deviation
  - Quantiles
- Histogram
  - Estimate of the probability distribution
  - pH outside the range 2.9–3.6 is uncommon

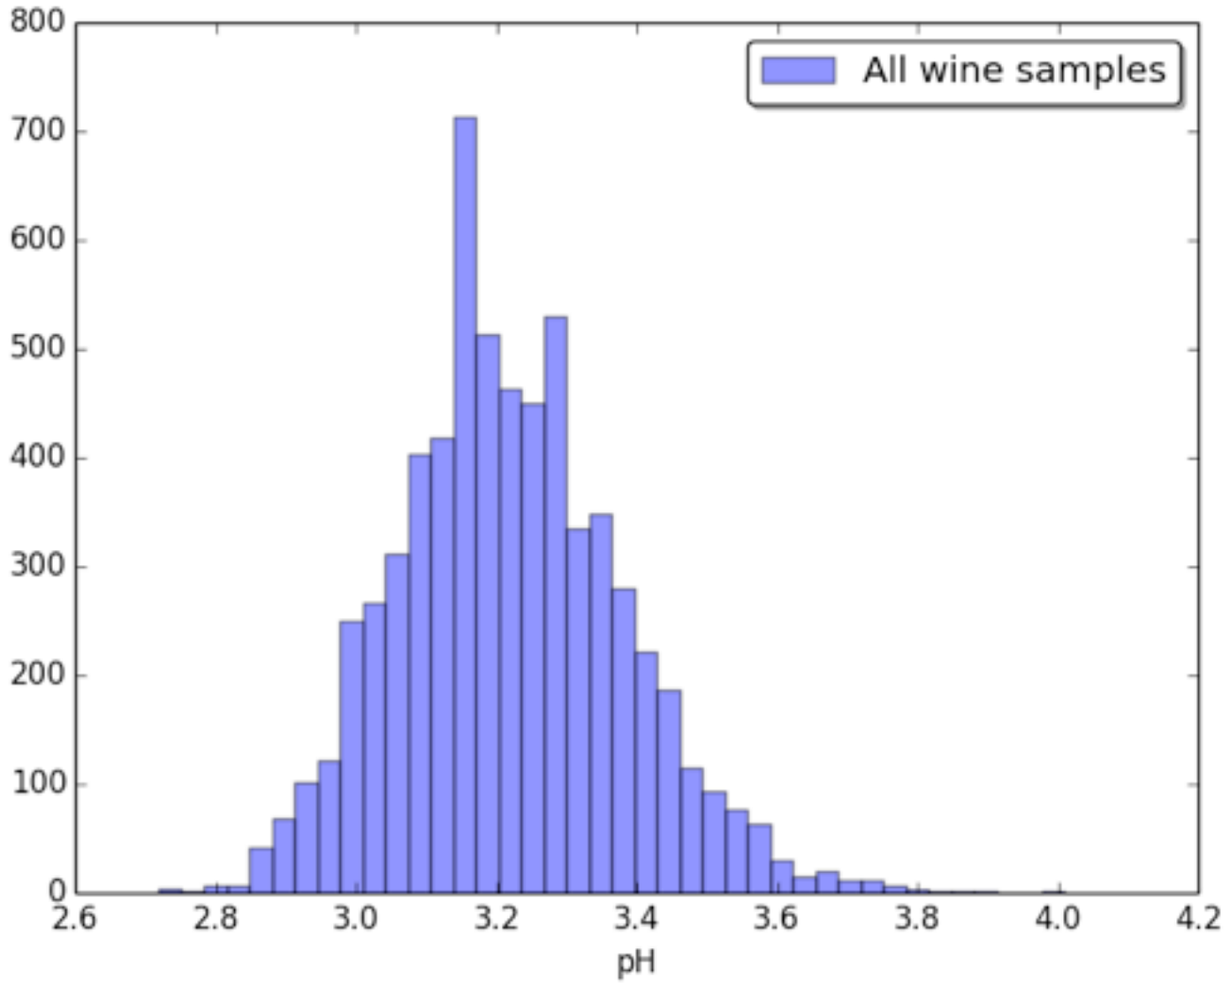




Reaktor

# Insights continued

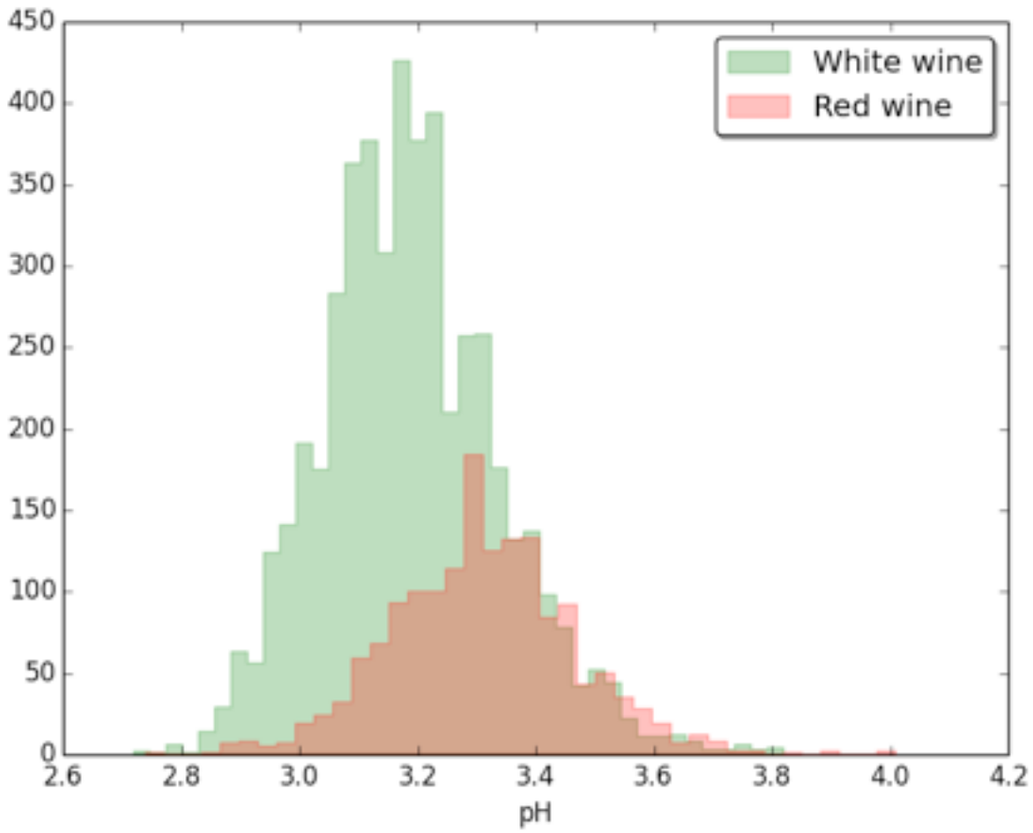


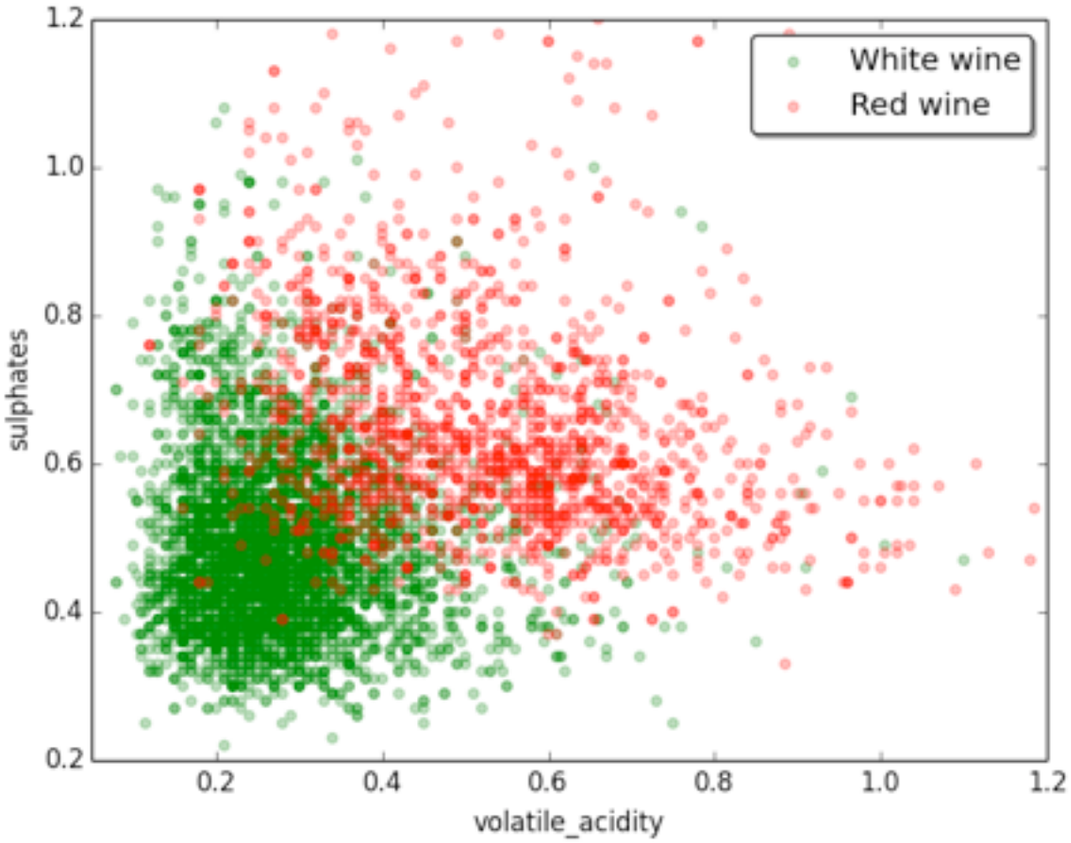


- Left: Histograms separately for red and white wines (1D)
- Right: A scatter plot
  - Not only variation but also covariation of two variables (2D)
- How about visualization in higher dimensions?



Reaktor

# Hands on

- **The first session includes:**
  - Read data from the file
  - Plot histograms of all variables
  - Plot histograms by wine type
  - Select two variables
  - Create a scatter plot using selected variables
  - See the script `hands-on.py`

# Modeling: Logistic regression

- **Objective**: Find decision boundary to classify new samples
- To create a model:
  - Select a model class
  - Fit the model: Maximize the probability of the observed data (maximum likelihood)
- **Output**:
  - Probability of measurements representing a red wine
  - If probability > 0.5, the sample is classified as red wine

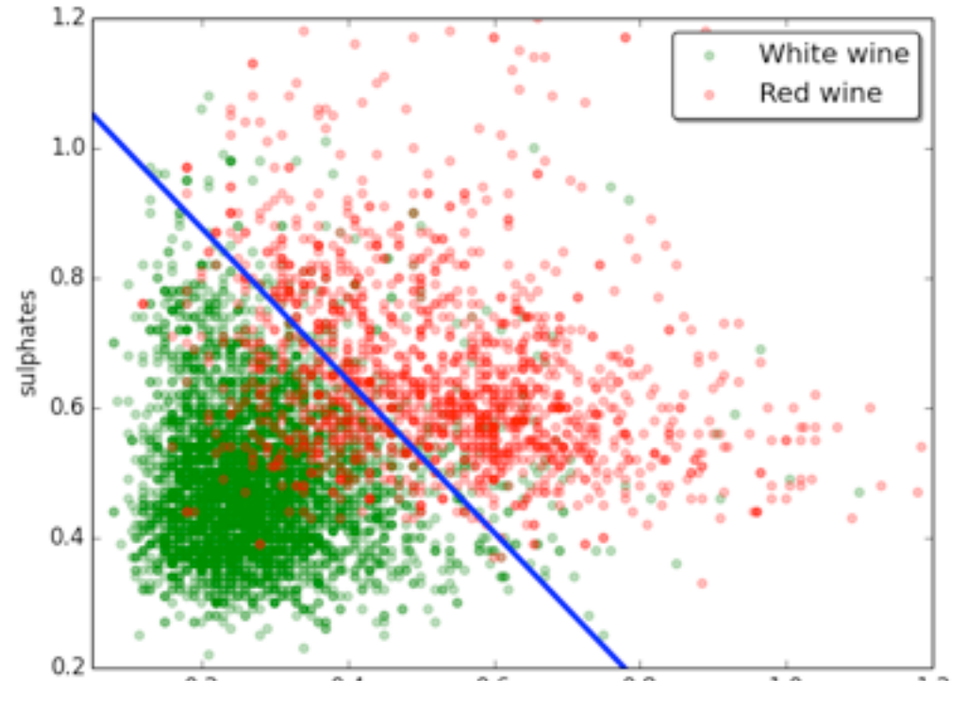


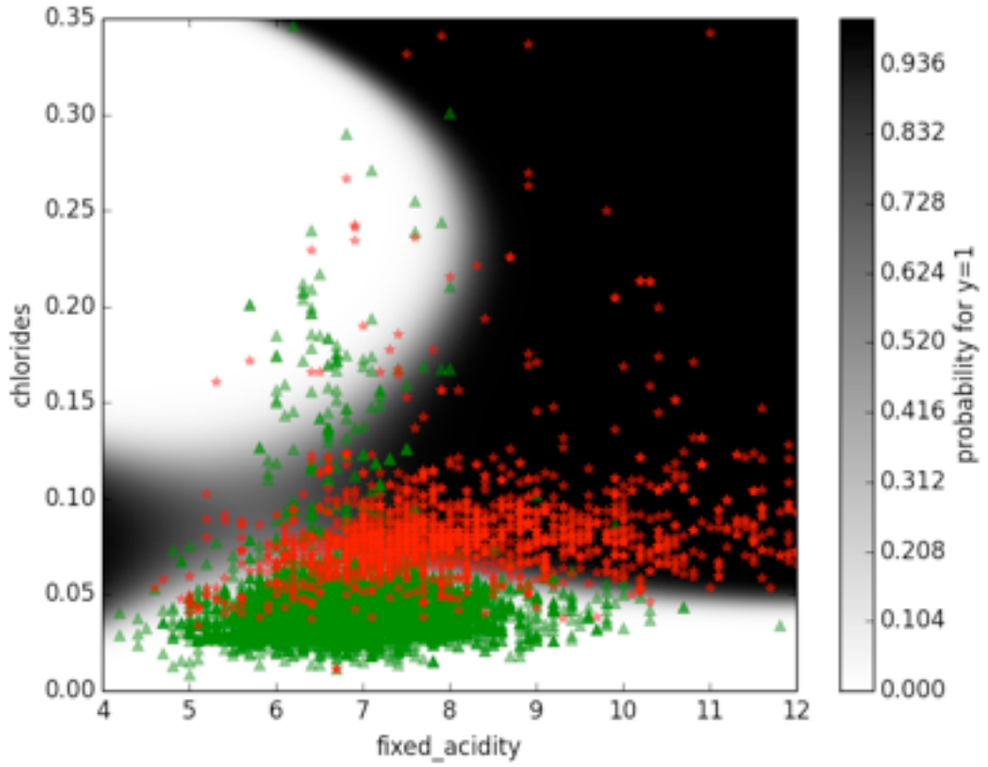




Reaktor

# Logistic regression

The conditional probabilities depend on the linear combination of variables. By inverse logit (ilogit) the outputs are limited to the range [0,1].

$$p(y=1|x,w) = \text{ilogit}(x^T w)$$

where $\text{ilogit}(u) = \dfrac{1}{1+e^{-u}}$

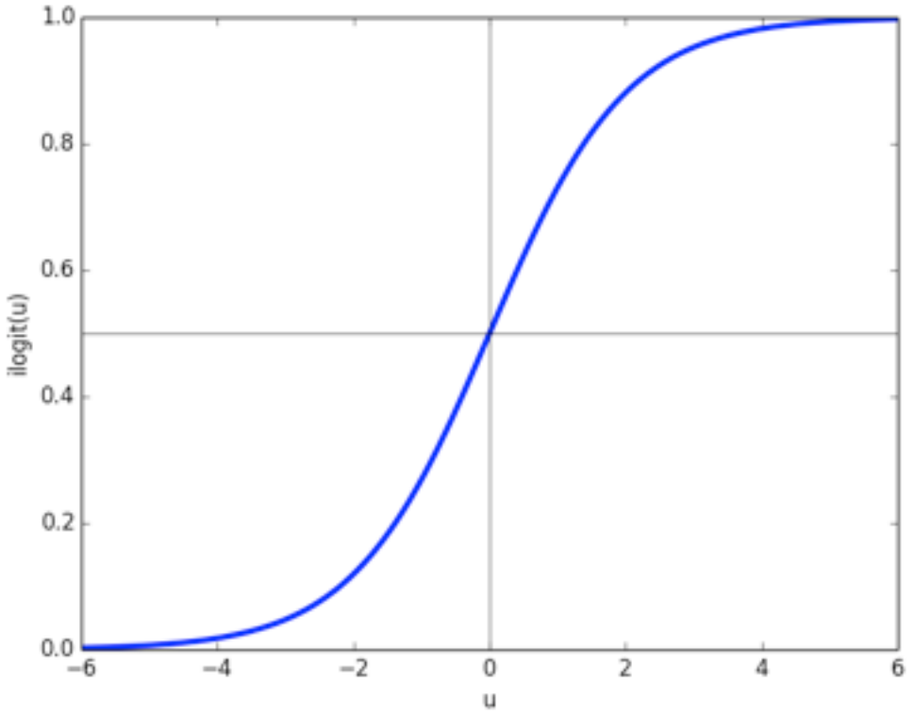


The model parameters are optimized by maximizing the probability of observed data. The maximum likelihood cost function is

$$\max_w \prod_{i=1}^{N} p_i^{y_i} (1-p_i)^{1-y_i}$$

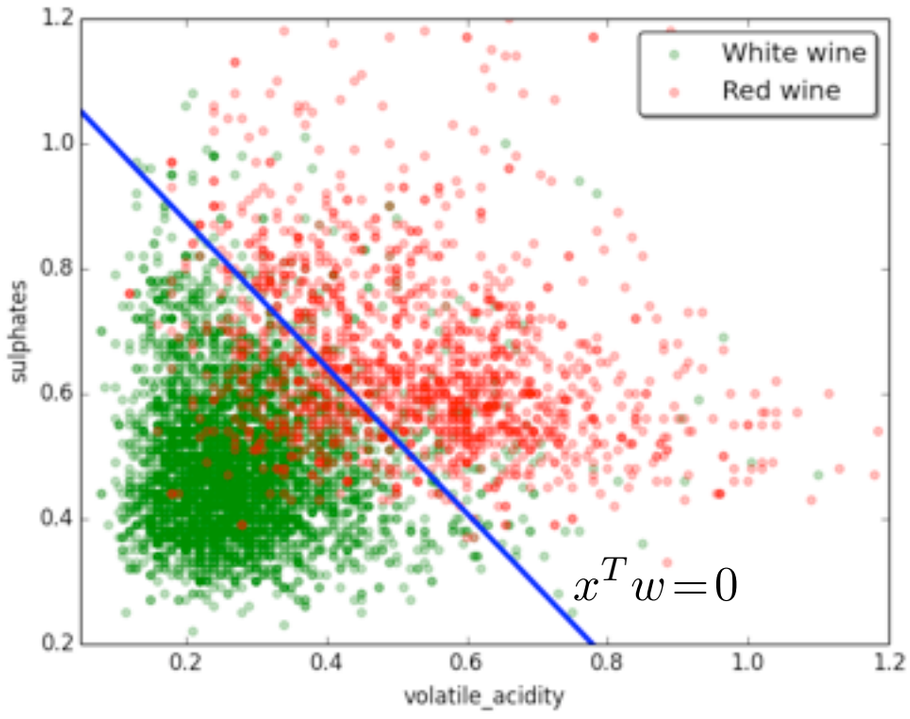




Reaktor

# Hands-on continued

- **The second session includes:**
  - Continue with the selected two variables
  - 1st model: A linear decision boundary
    - Visualize the result
  - 2nd model: A non-linear decision boundary
    - Create non-linearities using products and powers of the original variables
    - Visualize the result
  - See the script `hands-on.py`

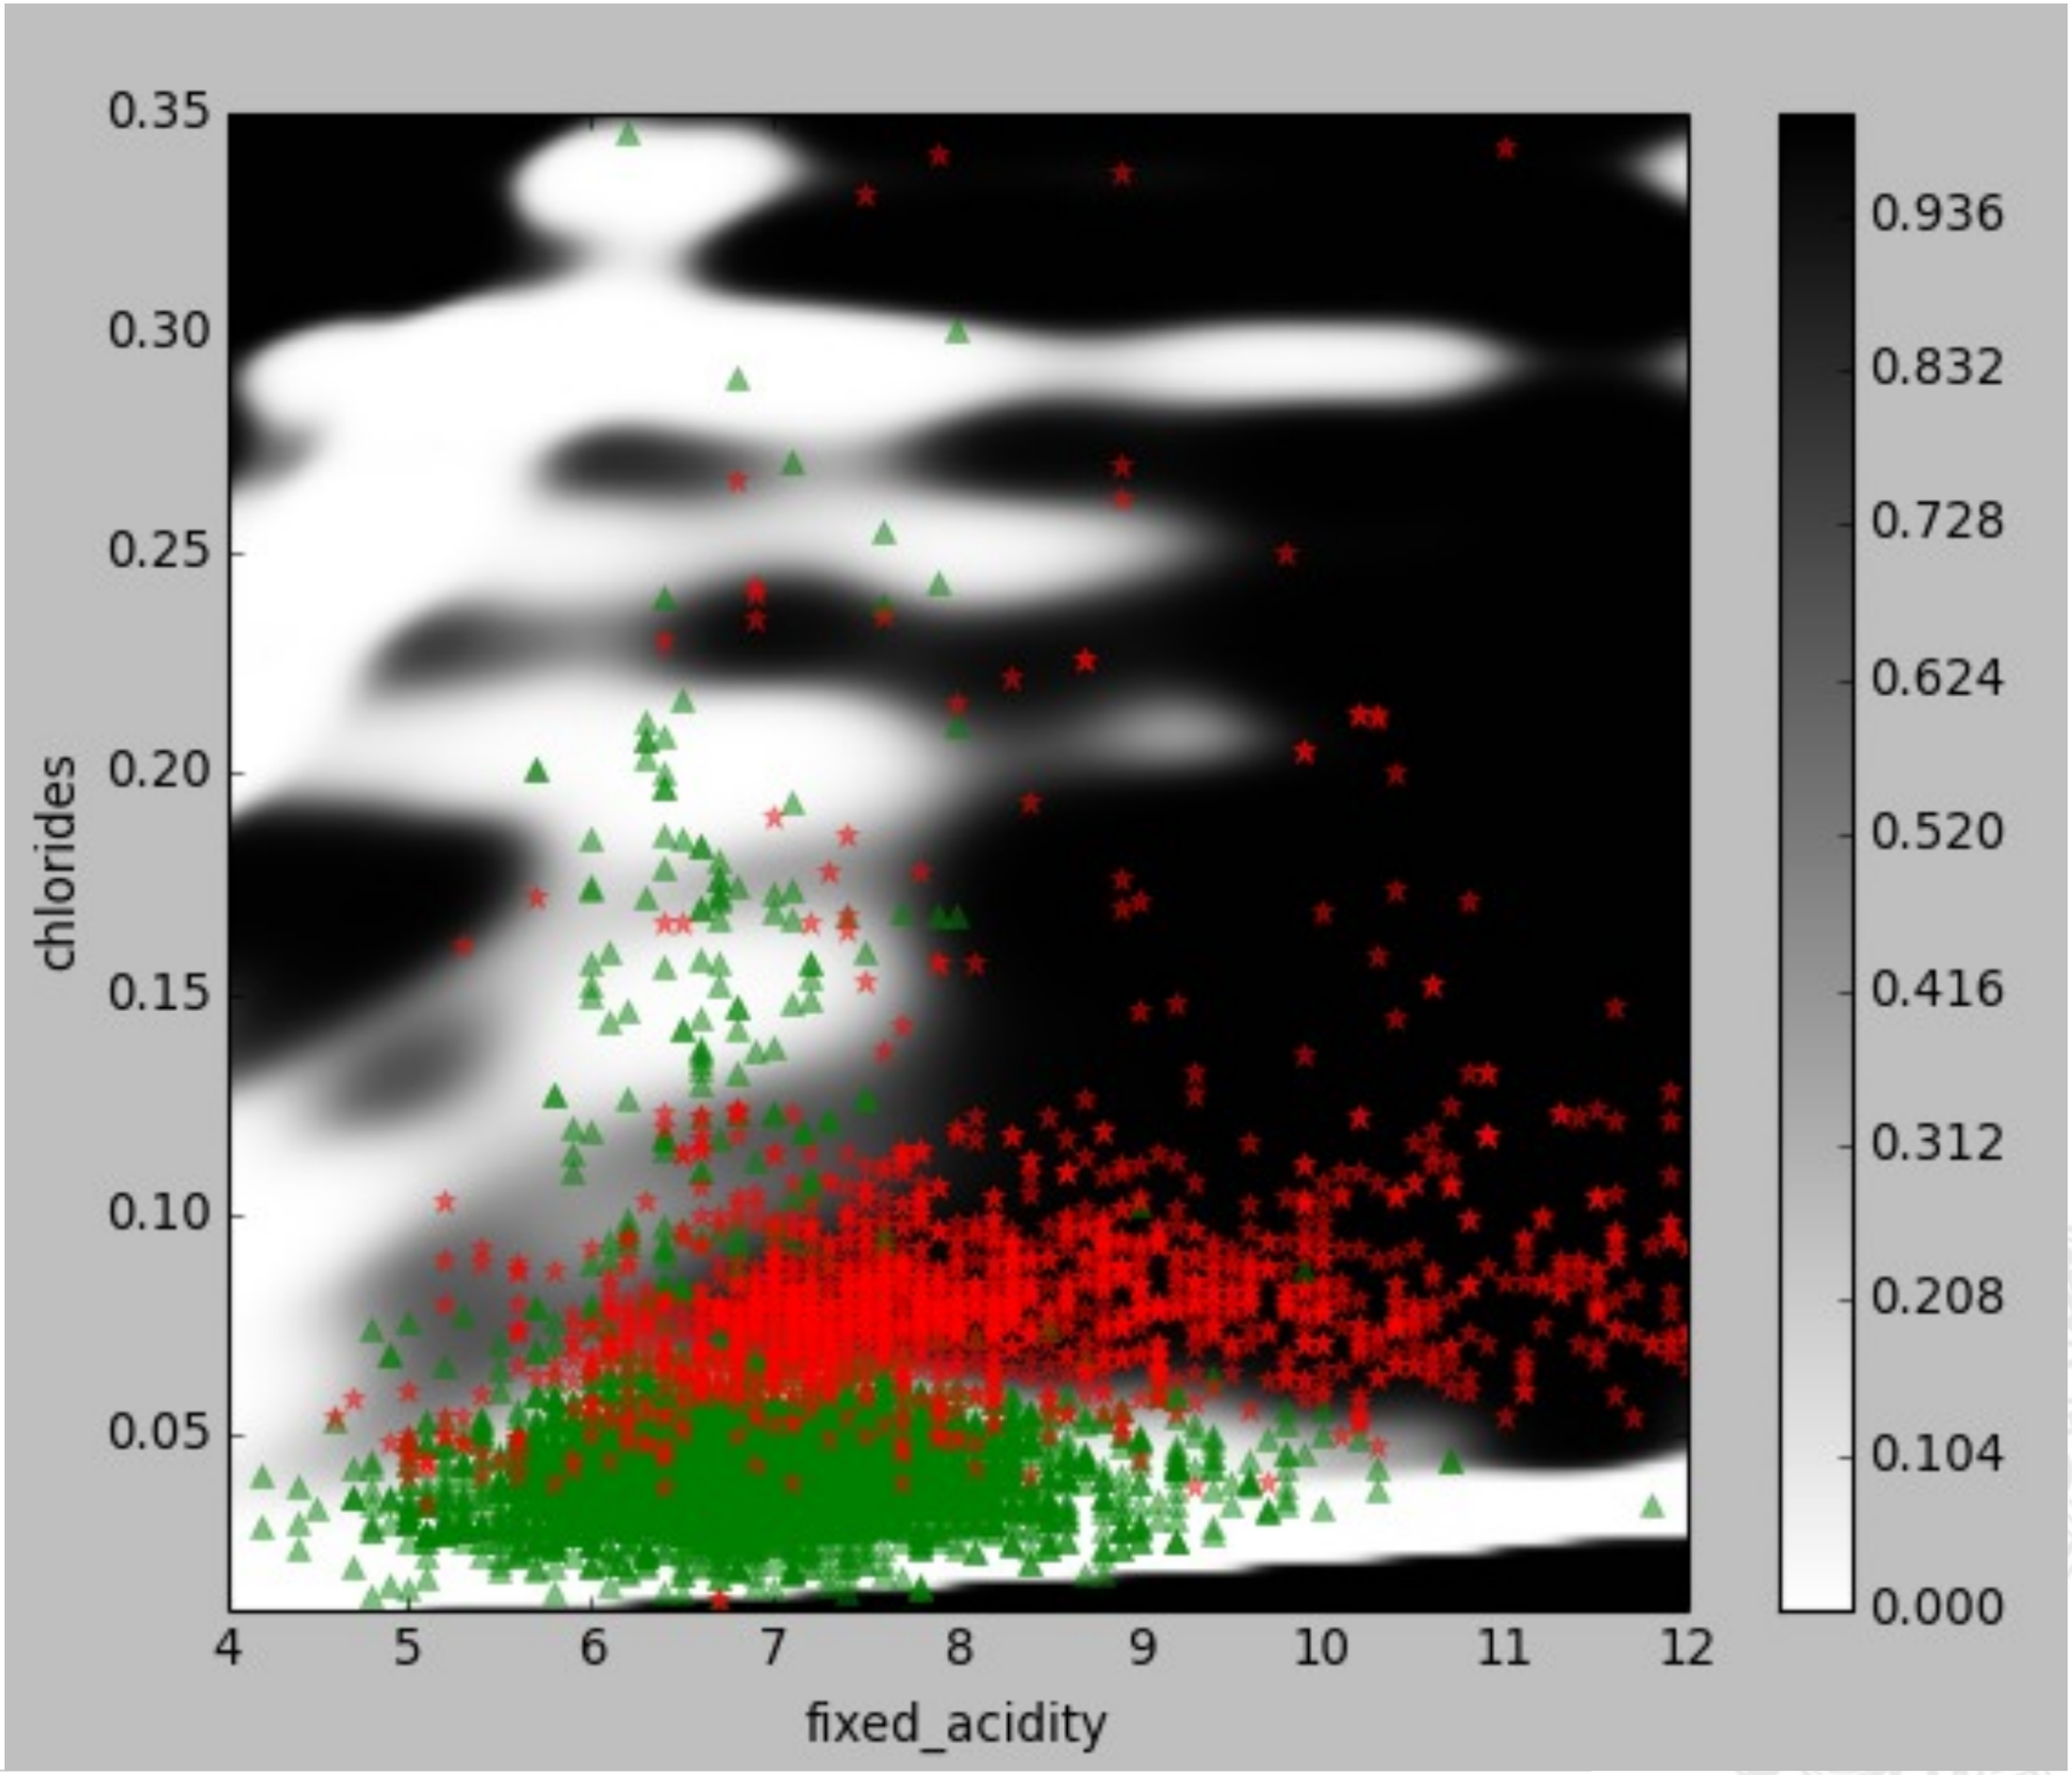




Reaktor

# Model selection continues...

- How to tell whether a model is good or not?
- The figure on the right visualizes a decision boundary. Is it a good classification rule?
- How to select a model out of numerous possibilities?
  - **The objective**: The model is supposed to *generalize* well, i.e. it should provide as accurate predictions as possible for *new* data (=data that was not used when fitting the model)

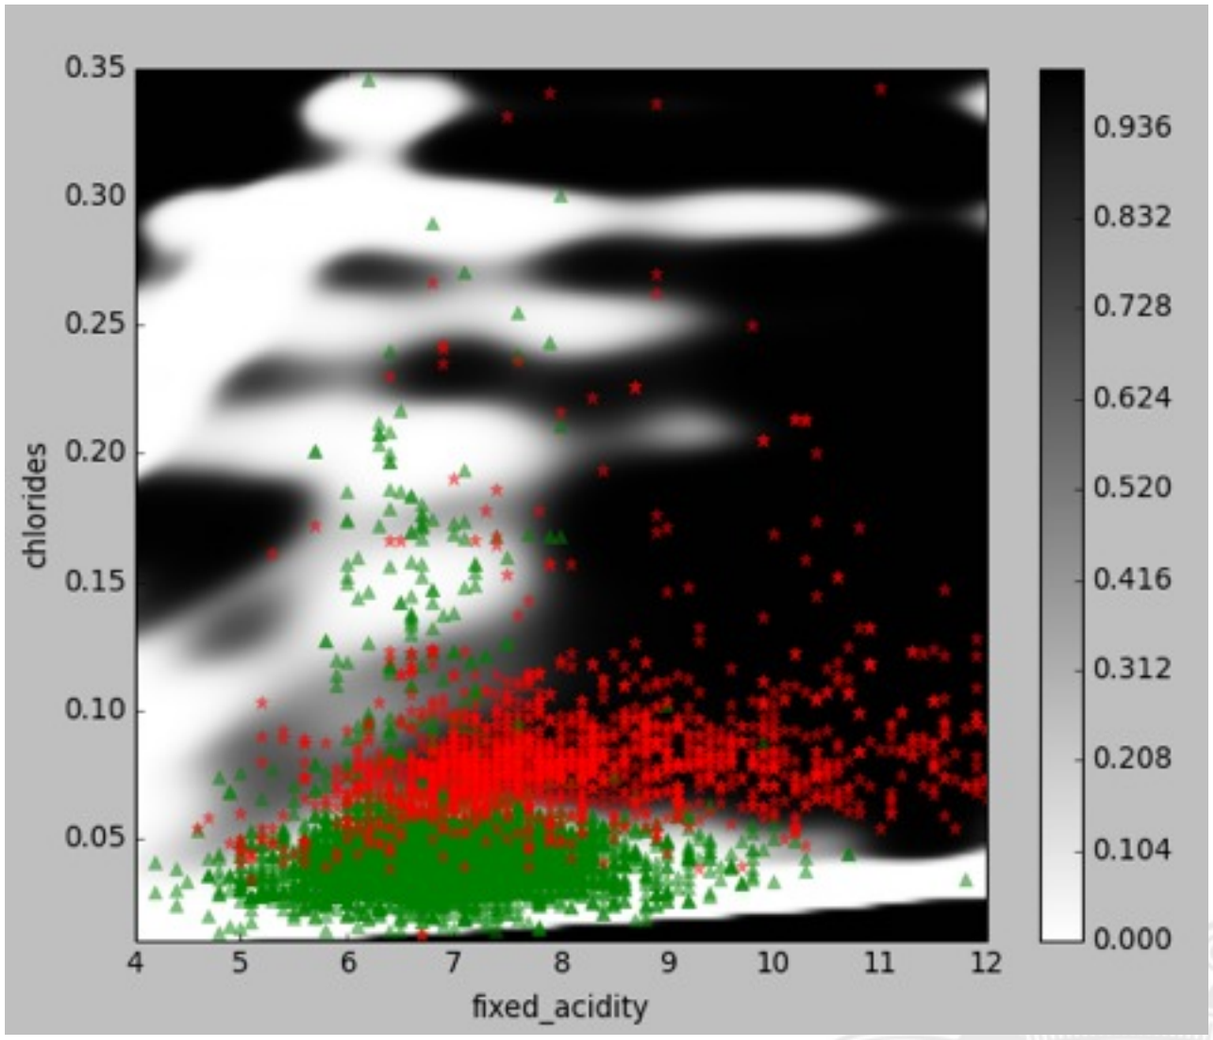


- Divide the data randomly to different parts:
  - **Training data set**: used to fit the model, i.e. optimize the parameters
  - **Validation/test data set**: used to evaluate prediction accuracy



Reaktor

# Hands on continued

- **The third session includes:**
  - Start by dividing the data set into two parts
  - Implement a process to test different models, e.g.
    - The best 2 variables out of available 11
    - With different nonlinear terms
    - The best K variables out of available 11
  - Find a model that generalizes well
    - Apply automated model selection
  - Hints in the end of the script `hands-on.py`

# Summary

- A classification problem was introduced
- The session covered
  - Data insights
  - Model fitting
  - Accuracy evaluation
  - Model selection
- In general, the process presented can be applied to any modeling task

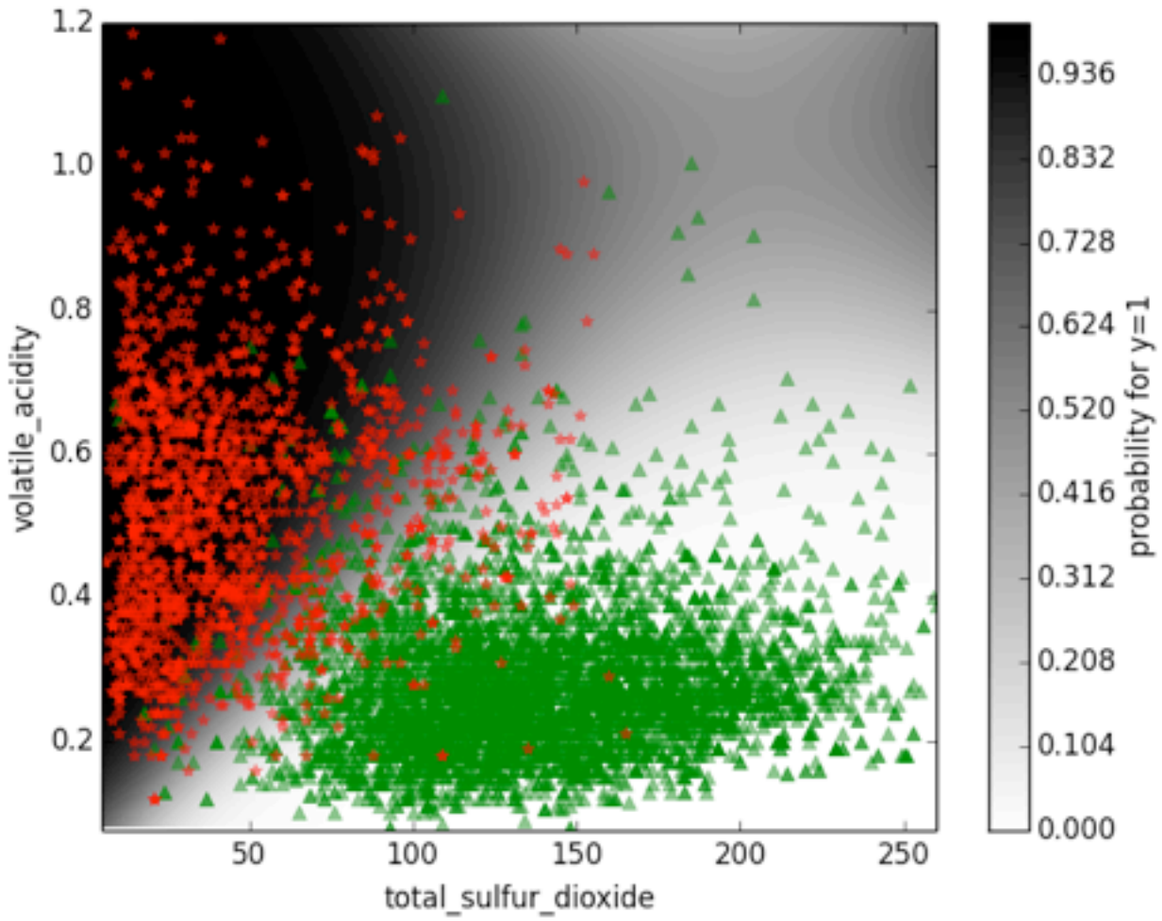


- **Two possible solutions**
  - Using two variables with quadratic transformations: test accuracy 95%
  - Using all the available variables and fitting with regularization
    - The model uses a linear combination of 8 variables with test accuracy 99.5%



Reaktor